

Technical Specification

This document outlines the technical requirements for the new data pipeline.

1. Overview

The pipeline ingests raw event data from multiple sources, transforms it, and loads it into the data warehouse on an hourly schedule.

2. Architecture

Components involved:

- Apache Kafka (ingestion)
- Apache Spark (transformation)
- Snowflake (storage)

3. SLAs

---	---	---	
Metric	Target	Current	
Latency	< 5 min	3.2 min	
Throughput	> 10K/s	12.4K/s	
Error rate	< 0.1%	0.04%	

4. Next Steps

Review and sign off on this spec by end of sprint.